



操作系统原理

Operating Systems Principles

张青
计算机学院



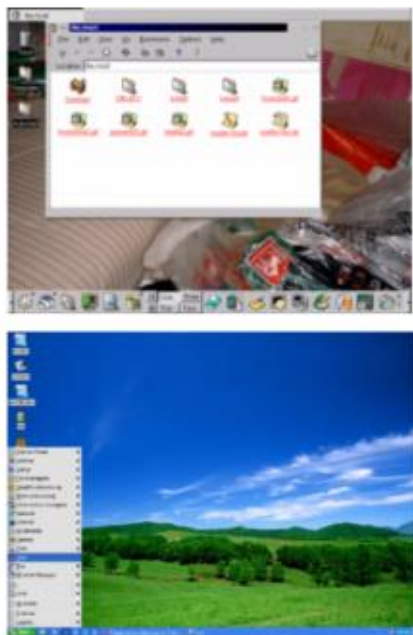
第三讲 — 中断、异常



基本概念和原理

中断 (hardware interrupt)

来自硬件设备（外设, device）的处理请求

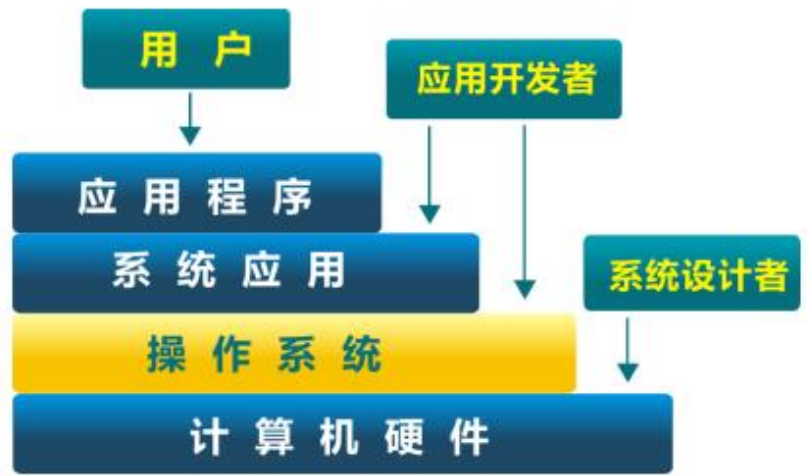



异步：产生原因和当前执行指令无关，如程序被磁盘读打断

基本概念和原理

异常 (exception)

非法指令或者其他原因导致**当前指令执行失败**, (如: 内存出错) 后的处理请求

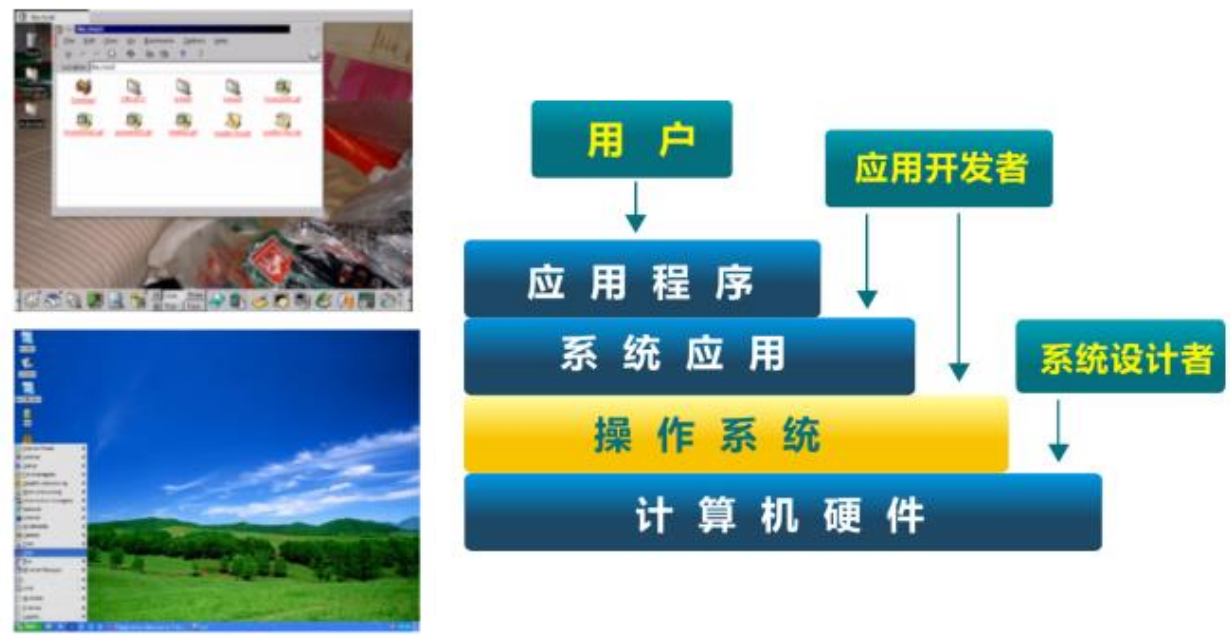


同步：产生和当前执行或试图执行的指令相关

基本概念和原理

系统调用 (system call)

应用程序**主动**向操作系统发出的服务请求



不同体系结构中的中断和异常

通用概念	产生原因	AArch64		x86-64
中断	硬件异步	异常	异步异常 (重置/中断)	中断 (可屏蔽/不可屏蔽)
异常	软件同步		同步异常 (终止/异常指令)	异常 (Fault/Trap/Abort)



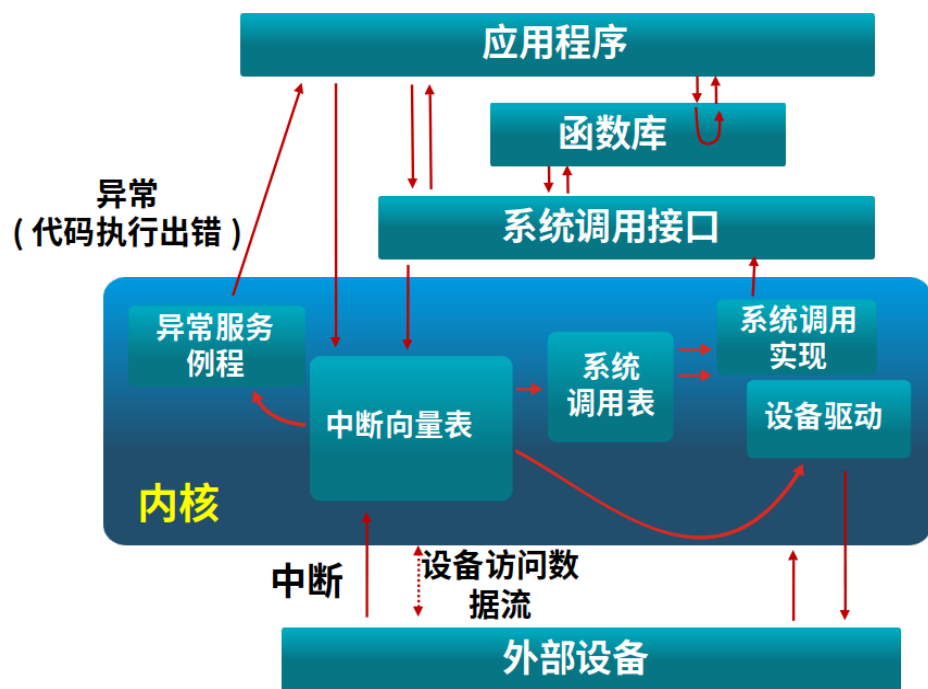
为什么需要



- 为什么需要中断、异常和系统调用
 - OS 内核是被信任的第三方
 - OS 内核可以执行特权指令，管理硬件
 - OS 内核提供了各种 Service
- 中断希望解决的问题
 - 当外设连接计算机时，会出现什么现象？
- 异常希望解决的问题
 - 当应用程序处理意想不到的行为时，会出现什么现象？
- 系统调用希望解决的问题
 - 用户应用程序是如何得到系统服务？



基本概念和原理



■ 源头

- 中断：外设
- 异常：应用程序意想不到的行为
- 系统调用：应用程序请求 OS 服务

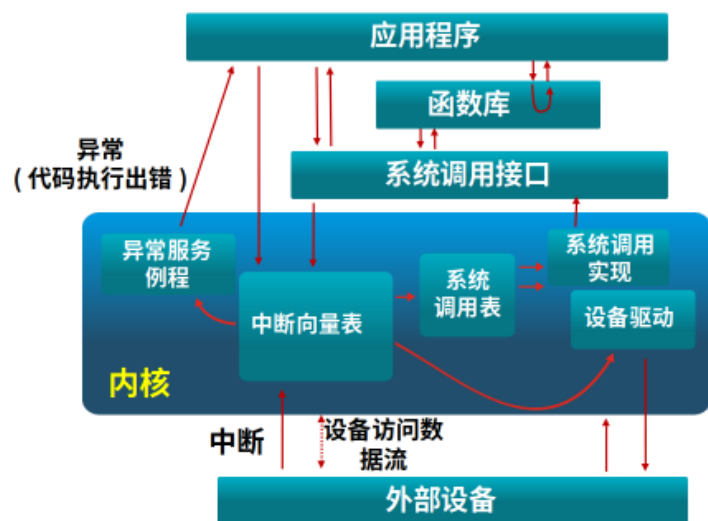
■ 响应方式

- 中断：异步
- 异常：同步
- 系统调用：异步或同步

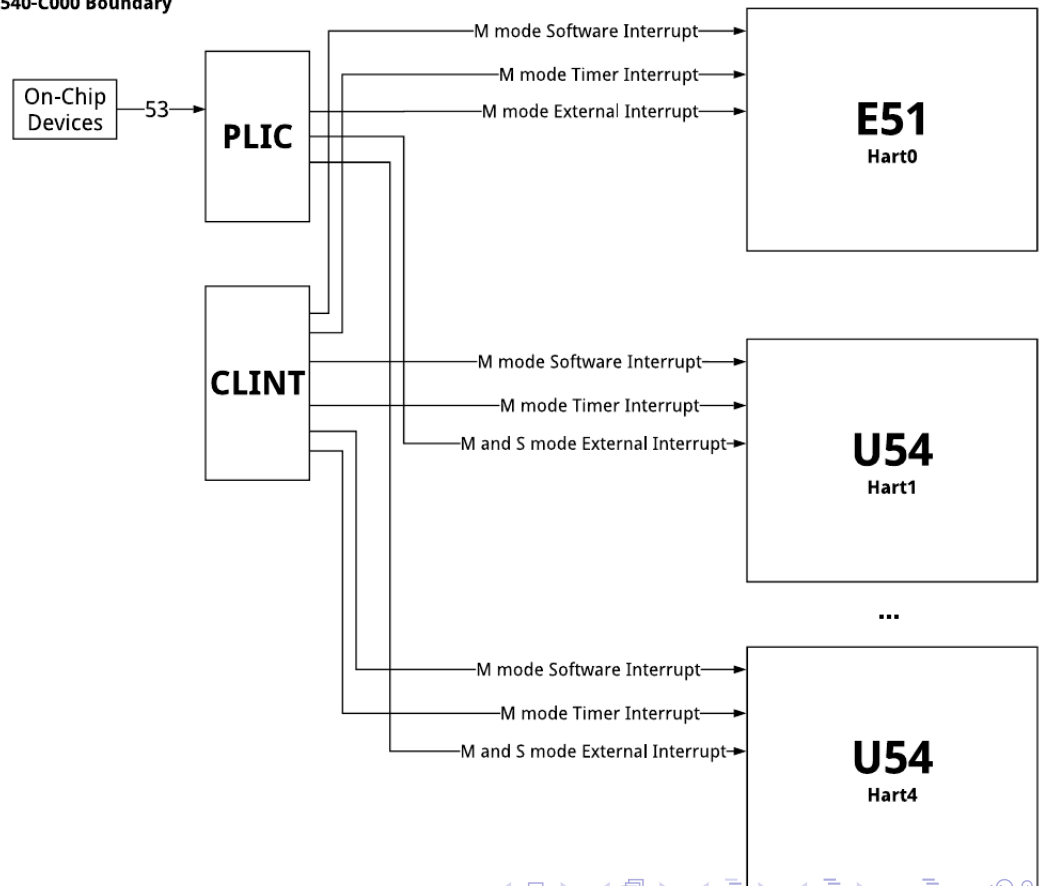
■ 处理机制

- 中断：持续，对应用程序透明
- 异常：杀死或者重新执行
- 系统调用：等待和持续

硬件支持



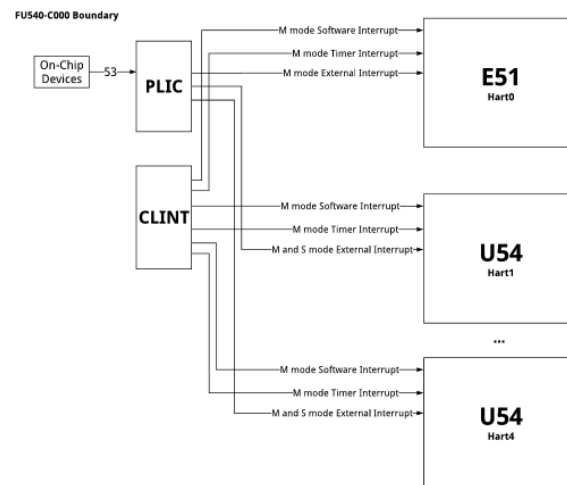
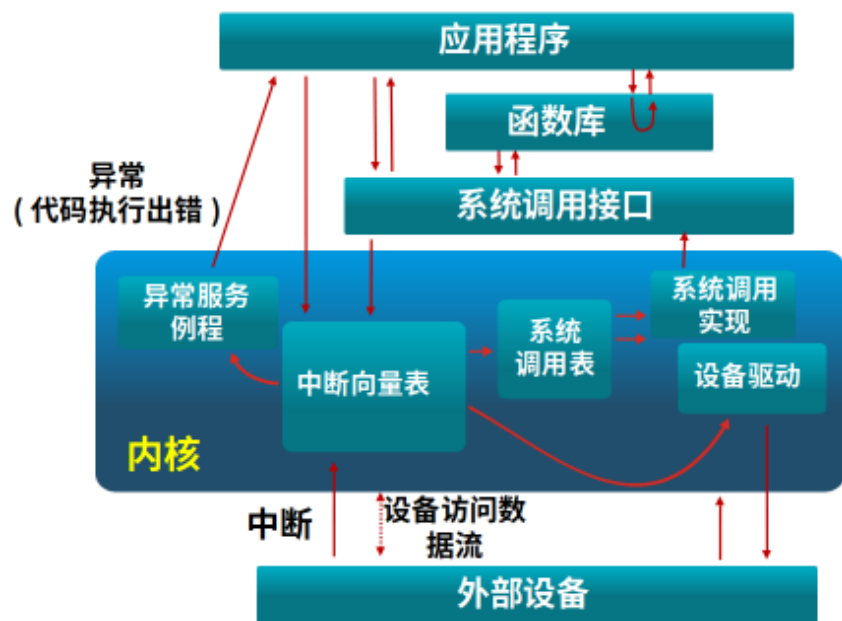
FU540-C000 Boundary



- Core Local Interruptor (CLINT)
- Platform-Level Interrupt Controller (PLIC)



硬件支持



三种标准的中断源:

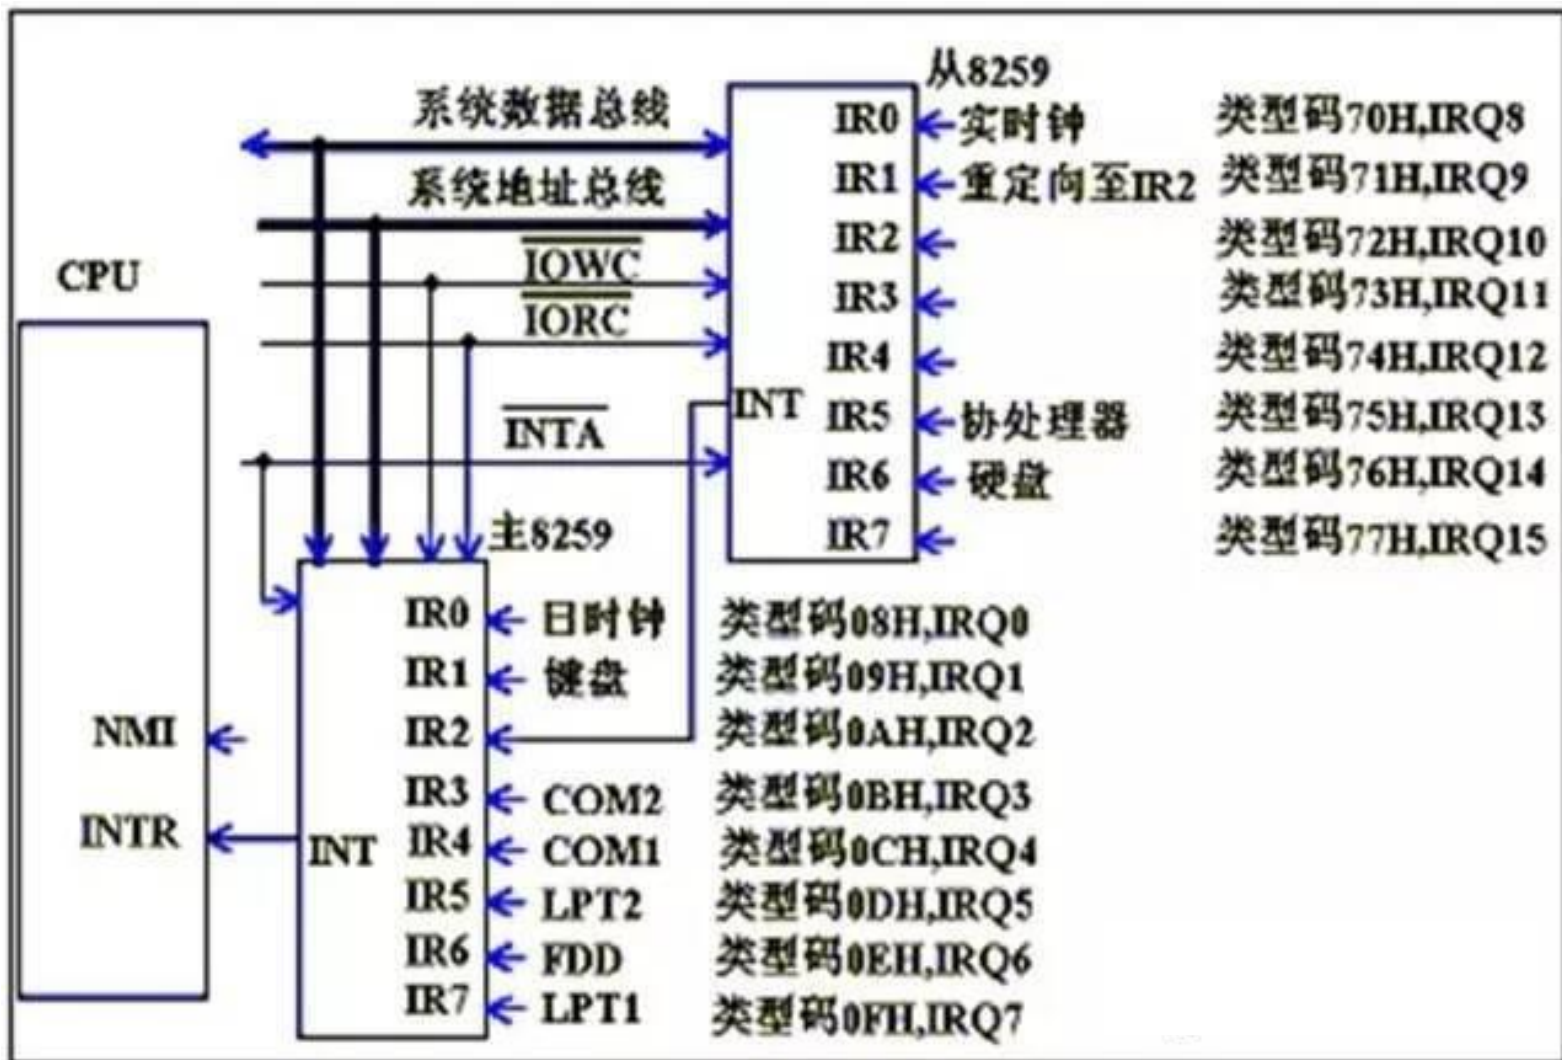
- Core Local Interruptor (CLINT)
- Platform-Level Interrupt Controller (PLIC)

- 软件中断通过向内存映射寄存器中存数来触发, 如 IPI
- 时钟中断, 如 $\text{stimecmp} > \text{stime}$
- 由平台级中断控制器引发外部中断



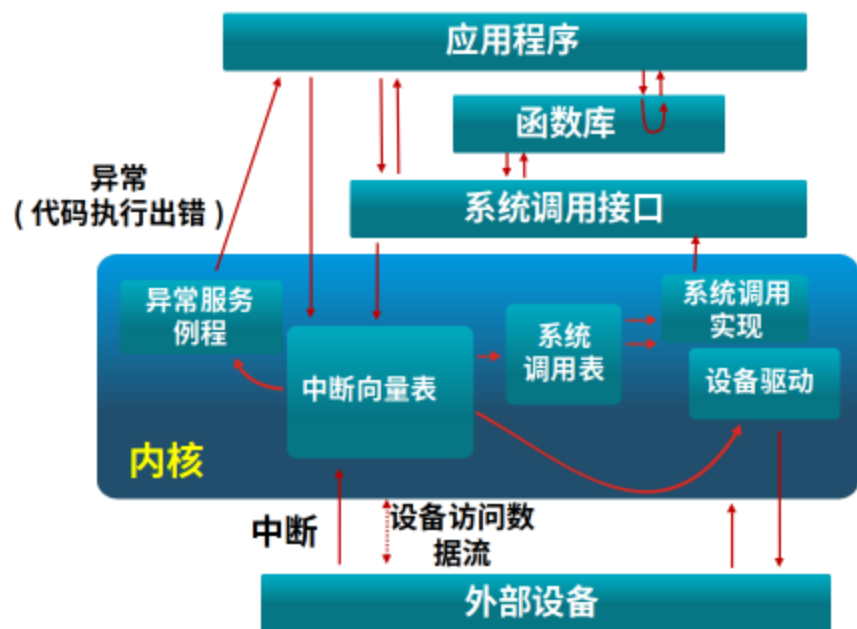
中断处理机制

X86平台 8259中断处理





中断处理机制



建立中断机制

- 建立中断服务例程
- 让 CPU 能响应中断
- 响应并处理中断
- 保存/恢复现场

- Core Local Interruptor (CLINT)
- Platform-Level Interrupt Controller (PLIC)



Linux系统中的中断

❖ 中断（设备产生、异步）

- 可屏蔽：设备产生的信号，通过中断控制器与处理器相连，可被暂时屏蔽（如，键盘、网络事件）
- 不可屏蔽：一些关键硬件的崩溃（如，内存校验错误）

❖ 异常（软件产生、同步）

- 错误（Fault）：如缺页异常（可恢复）、段错误（不可恢复）等
- 陷阱（Trap）：无需恢复，如断点（int 3）、系统调用（int 80）
- 中止（Abort）：严重的错误，不可恢复（机器检查）

Linux系统中的中断

- ❖ 在中断处理中做尽量少的事
- ❖ 推迟非关键行为
- ❖ 结构：Top half & Bottom half
 - Top half : 做最少的工作后返回
 - Bottom half : 推迟处理 (softirq, tasklets, 工作队列, 内核线程)





Top Half: 马上做

❖ 最小的、公共行为

- 保存寄存器、屏蔽其他中断
- 恢复寄存器，返回原来场景

❖ 最重要：调用合适的由硬件驱动提供的中断处理handler

❖ 因为中断被屏蔽，所以不要做太多事情（时间、空间）

❖ 使用将请求放入队列，或者设置标志位将其他处理推迟到 bottom half



Top Half: 找到handler

- ❖ 现代处理器中，多个I/O设备共享一个IRQ和中断向量
- ❖ 多个ISR (interrupt service routines) 可以结合在一个向量上
- ❖ 调用每个设备对应该IRQ的ISR



Bottom Half: 延迟完成

❖ 提供一些推迟完成任务的机制

- softirqs
- tasklets (建立在softirqs之上)
- 工作队列
- 内核线程

❖ 这些工作可以被中断



注意：中断处理没有进程上下文

- ❖ 中断（和异常相比）和具体的某条指令无关
- ❖ 也和中断时正在跑的进程、用户程序无关
- ❖ 中断处理handler不能睡眠！



中断处理中的一些约束条件

- ❖ 不能睡眠
 - 或者调用可能会睡眠的任务
- ❖ 不能调用`schedule()`调度
- ❖ 不能释放信号或调用可能睡眠的操作
- ❖ 不能和用户地址空间交换数据

Demo



如何安全执行中断

- **Interrupt Vector Table**

- Where the processor looks for a handler
- Limited number of entry points into kernel
- Stored in RAM at a known address

- **Atomic transfer of control**

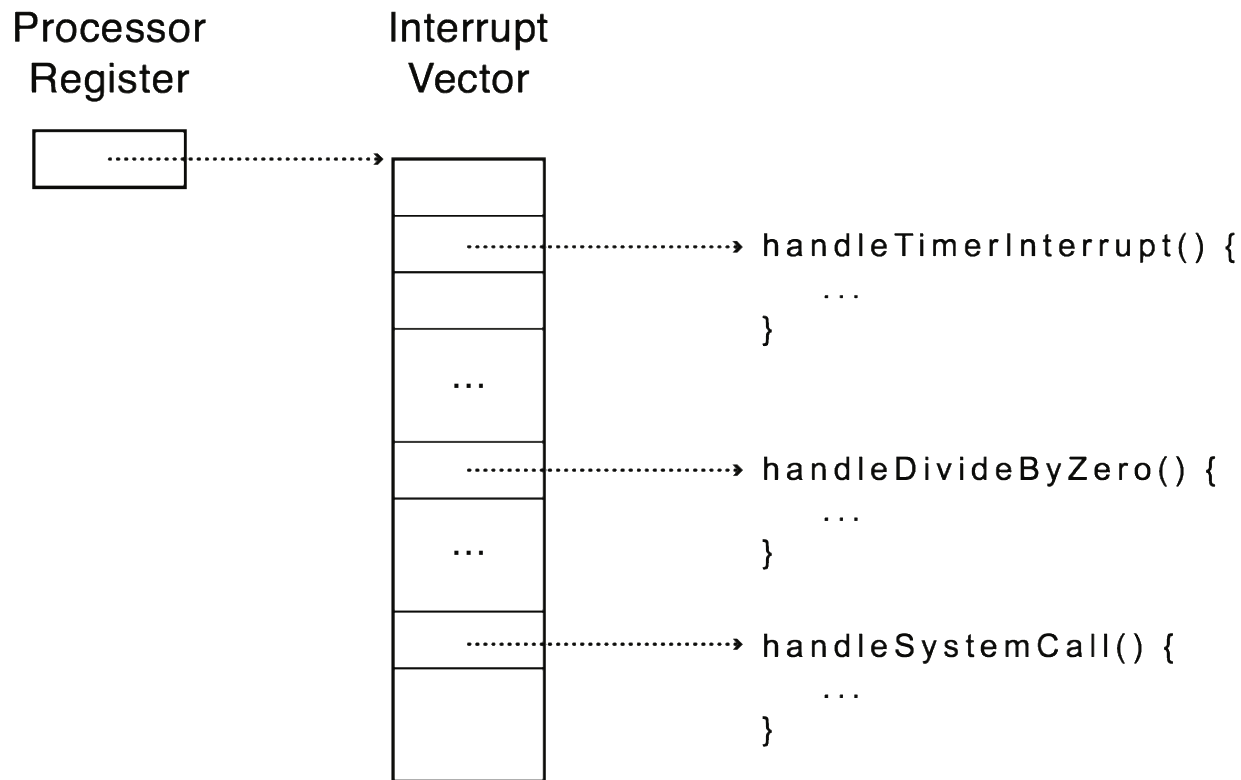
- Single instruction to change:
 - Program counter
 - Stack pointer
 - Memory protection
 - Kernel/user mode

- **Transparent restartable execution**

- User program does not know interrupt occurred

中断向量表

Table set up by OS kernel; pointers to code to run on different events

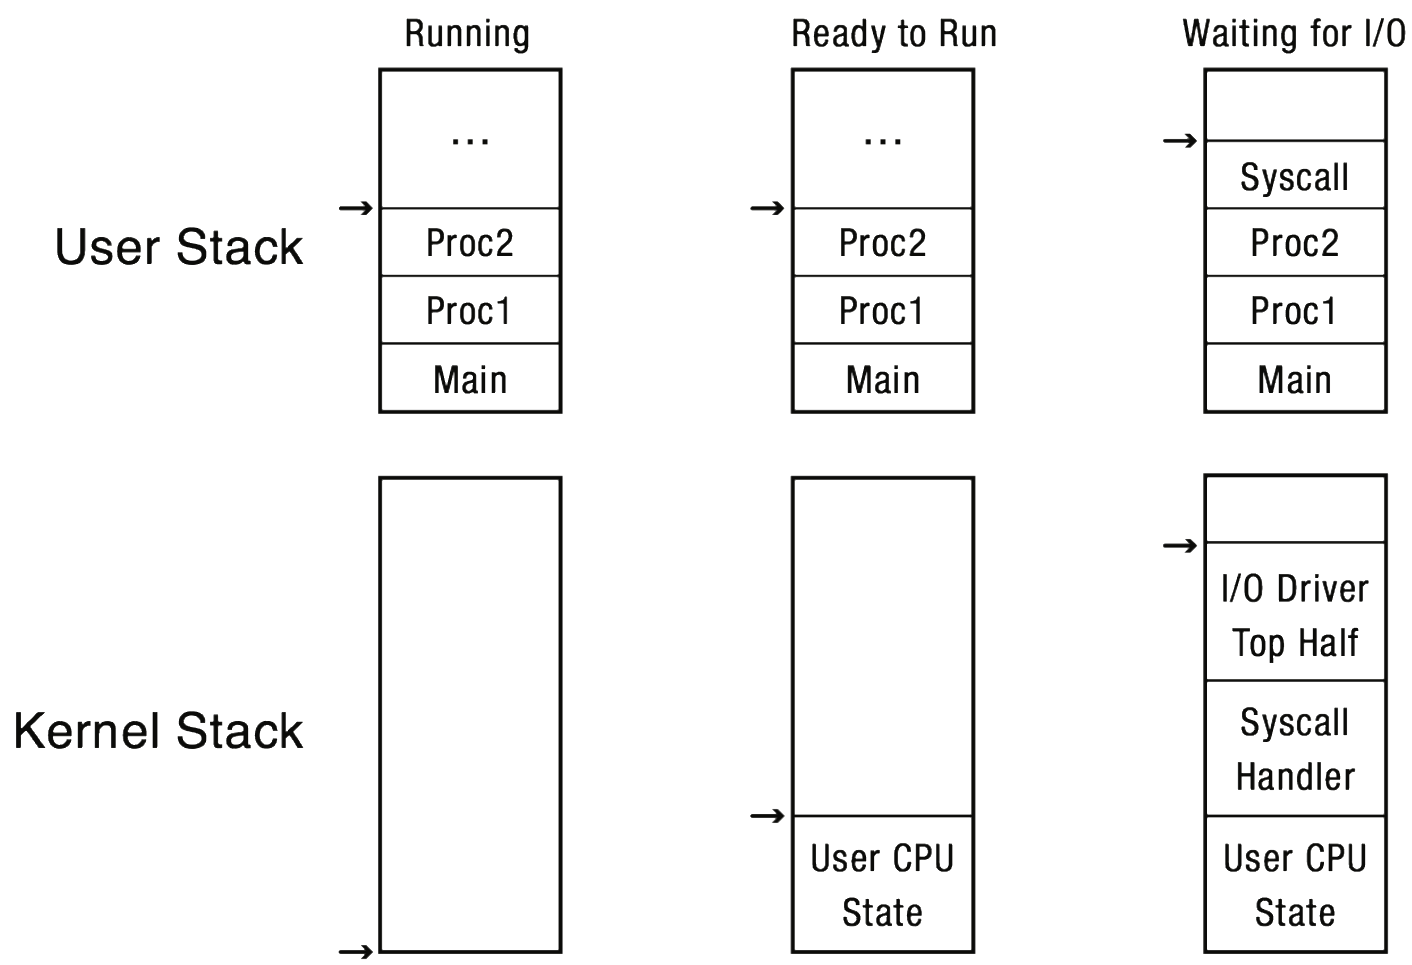




中断向量表

- Per-processor, located in kernel (not user) memory
 - Fun fact! Usually a process/thread has both a kernel and user stack
- **Can the interrupt handler run on the stack of the interrupted user process?**

中断向量表



为什么使用INT

- Hardware devices may need asynchronous and immediate service. For example:
 - Timer interrupt: Timers and time-dependent activities need to be updated with the passage of time at precise intervals
 - Network interrupt: The network card interrupts the CPU when data arrives from the network
 - I/O device interrupt: I/O devices (such as mouse and keyboard) issue hardware interrupts when they have input (e.g., a new character or mouse click)

多核上的中断

- How are interrupts handled on multicore machines?
 - On x86 systems each CPU gets its own local Advanced Programmable Interrupt Controller (APIC). They are wired in a way that allows routing device interrupts to any selected local APIC.
 - The OS can program the APICs to determine which interrupts get routed to which CPUs.
 - The default (unless OS states otherwise) is to route all interrupts to processor 0

其他类型的中断

- Software Interrupts:
 - Interrupts caused by the execution of a software instruction:
 - `INT <interrupt_number>`
 - Used by the system call `interrupt()`
- Initiated by the running (user level) process
- Cause current processing to be interrupted and transfers control to the corresponding interrupt handler in the kernel



其他类型的中断

- Exceptions
 - Initiated by processor hardware itself
 - Example: divide by zero
- Like a software interrupt, they cause a transfer of control to the kernel to handle the exception



其他类型的中断

- Exceptions
 - Initiated by processor hardware itself
 - Example: divide by zero
- Like a software interrupt, they cause a transfer of control to the kernel to handle the exception



其他类型的中断

- HW -> CPU -> Kernel: Classic HW Interrupt
- User -> Kernel: SW Interrupt
- CPU -> Kernel: Exception
- Interrupt Handlers used in all 3 scenarios

中断屏蔽

- Interrupt handler runs with interrupts off
 - Re-enabled when interrupt completes
- OS kernel can also turn interrupts off
 - Eg., when determining the next process/thread to run

中断处理程序

Designing an Interrupt Handler:

- Since the interrupt handler must be minimal, all other processing related to the event that caused the interrupt must be deferred
 - Example:
 - Network interrupt causes packet to be copied from network card
 - Other processing on the packet should be deferred until its time comes
- The deferred portion of interrupt processing is called the “Bottom Half”

Bottom half

- Method for deferring portion of interrupt processing
- Globally serialized
 - When one bottom half is executing, no other bottom half can execute (even different type) on any CPU.
- Obvious performance limitations; primarily available for legacy support.
- Note: other mechanisms for deferred work are also sometimes referred to as bottom half mechanisms.

软中断

- Handlers that, like bottom halves, must be statically defined/allocated in the Linux kernel at compile time.
- A hardware interrupt handler (before returning) uses `raise_softirq()` to mark that a given `soft_irq` must execute deferred work
- At a later time, when scheduling permits, the marked `soft_irq` handler is executed
 - When a hardware interrupt is finished
 - When a process makes a system call
 - When a new process is scheduled
- Unlike bottom halves, softirqs are reentrant and can be executed concurrently on several CPUs

谢谢